

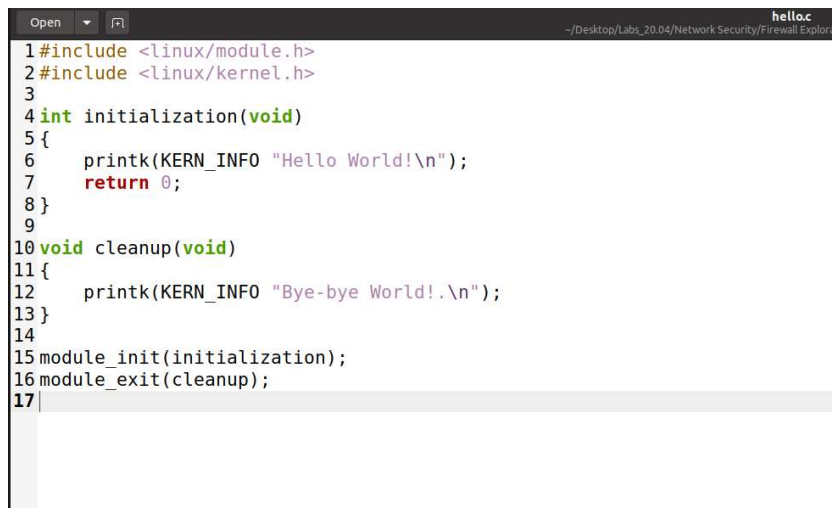
# Firewall Exploration Lab

57118210 郑荔文

## 3 Task 1: Implementing a Simple Firewall

### 3.1 Task 1.A: Implement a Simple Kernel Module

在 file 中有如下 hello.c 文件，其功能为在模块加载时打印“Hello World! ”，在模块被移除时打印“Bye-bye World! ”并且这些信息在 dmesg 命令下可以查看

A screenshot of a code editor window titled 'hello.c' with a file path of '~/.Desktop/Labs\_20.04/Network Security/Firewall Exploration'. The code is as follows:

```
1#include <linux/module.h>
2#include <linux/kernel.h>
3
4int initialization(void)
5{
6    printk(KERN_INFO "Hello World!\n");
7    return 0;
8}
9
10void cleanup(void)
11{
12    printk(KERN_INFO "Bye-bye World!\n");
13}
14
15module_init(initialization);
16module_exit(cleanup);
17
```

在同目录下还有一个 Makefile 文件，可以将上述 hello.c 文件编译成为一个可加载的内核模块

A screenshot of a terminal window titled 'Terminal'. The content of the Makefile is as follows:

```
obj-m += hello.o

all:
    make -C /lib/modules/$(shell uname -r)/build M=$(shell pwd) modules

clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean

~
~
~
```

之后在 `sudo make` 进行编译，再采用 `sudo insmod hello.ko` 插入模块，采用 `lsmod | grep hello` 列出模块信息，`sudo rmmod hello` 移除该模块，采用 `dmesg` 查看是否出现了“Hello World! ”和“Bye-bye World! ”信息

```

[07/23/21]seed@VM:~/.../kernel_module$ sudo make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/Desktop/Labs_20.04/Network_Security/Firewall_Exploration_Lab/Labsetup/Files/kernel_module modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M] /home/seed/Desktop/Labs_20.04/Network_Security/Firewall_Exploration_Lab/Labsetup/Files/kernel_module/hello.o
  Building modules, stage 2.
  MODPOST 1 modules
WARNING: modpost: missing MODULE_LICENSE() in /home/seed/Desktop/Labs_20.04/Network_Security/Firewall_Exploration_Lab/Labsetup/Files/kernel_module/hello.o
see include/linux/module.h for more information
  CC [M] /home/seed/Desktop/Labs_20.04/Network_Security/Firewall_Exploration_Lab/Labsetup/Files/kernel_module/hello.mod.o
  LD [M] /home/seed/Desktop/Labs_20.04/Network_Security/Firewall_Exploration_Lab/Labsetup/Files/kernel_module/hello.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
[07/23/21]seed@VM:~/.../kernel_module$ sudo insmod hello.ko
sudo: insmod: command not found
[07/23/21]seed@VM:~/.../kernel_module$ sudo insmod hello.ko
[07/23/21]seed@VM:~/.../kernel_module$ lsmod | grep hello
hello                16384  0
[07/23/21]seed@VM:~/.../kernel_module$ sudo rmmod hello

[07/23/21]seed@VM:~/.../kernel_module$ dmesg
[  0.000000] Linux version 5.4.0-54-generic (buildd@lcy01-amd64-024) (gcc version 9.3.0 (Ubuntu 9.3.0-17ubuntu1~20.04)) #60-Ubuntu SMP Fri Nov 6 10:37:59 UTC 2020 (Ubuntu 5.4.0-54.60-generic 5.4.65)

[ 340.703735] hello: loading out-of-tree module taints kernel.
[ 340.703737] hello: module license 'unspecified' taints kernel.
[ 340.703737] Disabling lock debugging due to kernel taint
[ 340.703757] hello: module verification failed: signature and/or required key missing - tainting kernel
[ 340.704325] Hello World!
[ 374.861114] Bye-bye World!.

```

可见再 dmesg 中出现了“Hello World! ”和“Bye-bye World! ”，与预期结果相符，表明该模块被插入内核模块中，之后又被移除。

## 3.2 Task 1.B: Implement a Simple Firewall Using Netfilter

### Tasks

1.使用 Makefile 编译 seedFilter.c,并将其加载入内核，使用 dig @8.8.8.8 www.example .com 命令产生 UDP 包，此时申请被阻塞，得不到应答

```

[07/23/21]seed@VM:~/.../packet_filter$ sudo insmod seedFilter.ko
[07/23/21]seed@VM:~/.../packet_filter$ lsmod | grep seedFilter
seedFilter            16384  0
[07/23/21]seed@VM:~/.../packet_filter$ dig @8.8.8.8 www.example.com

```

```

; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; connection timed out; no servers could be reached

```

将该模块移除后，重复上述命令，则可以得到应答，说明防火墙起效

```
[07/23/21]seed@VM:~/.../kernel_module$ dig @8.8.8.8 www.example.com

; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 43581
;; flags: qr rd ra ad; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags::; udp: 512
;; QUESTION SECTION:
;www.example.com.                IN      A

;; ANSWER SECTION:
www.example.com.                21265   IN      A      93.184.216.34

;; Query time: 259 msec
;; SERVER: 8.8.8.8#53(8.8.8.8)
;; WHEN: Fri Jul 23 07:09:14 EDT 2021
;; MSG SIZE rcvd: 60
```

## 2. 在 seedFilter.c 中修改相关代码，增加 hook

```
74 int registerFilter(void) {
75     printk(KERN_INFO "Registering filters.\n");
76
77     hook1.hook = printInfo;
78     hook1.hooknum = NF_INET_LOCAL_OUT;
79     hook1.pf = PF_INET;
80     hook1.priority = NF_IP_PRI_FIRST;
81     nf_register_net_hook(&init_net, &hook1);
82
83     hook3.hook = printInfo;
84     hook3.hooknum = NF_INET_LOCAL_IN;
85     hook3.pf = PF_INET;
86     hook3.priority = NF_IP_PRI_FIRST;
87     nf_register_net_hook(&init_net, &hook3);
88
89     hook4.hook = printInfo;
90     hook4.hooknum = NF_INET_FORWARD;
91     hook4.pf = PF_INET;
92     hook4.priority = NF_IP_PRI_FIRST;
93     nf_register_net_hook(&init_net, &hook4);
94
95     hook5.hook = printInfo;
96     hook5.hooknum = NF_INET_PRE_ROUTING;
97     hook5.pf = PF_INET;
98     hook5.priority = NF_IP_PRI_FIRST;
99     nf_register_net_hook(&init_net, &hook5);
100
101     hook6.hook = printInfo;
102     hook6.hooknum = NF_INET_POST_ROUTING;
103     hook6.pf = PF_INET;
104     hook6.priority = NF_IP_PRI_FIRST;
105     nf_register_net_hook(&init_net, &hook6);
106
107     hook2.hook = blockUDP;
108     hook2.hooknum = NF_INET_POST_ROUTING;
109     hook2.pf = PF_INET;
110     hook2.priority = NF_IP_PRI_FIRST;
111     nf_register_net_hook(&init_net, &hook2);
112
113     return 0;
114 }
115
116 void removeFilter(void) {
117     printk(KERN_INFO "The filters are being removed.\n");
118     nf_unregister_net_hook(&init_net, &hook1);
119     nf_unregister_net_hook(&init_net, &hook2);
120     nf_unregister_net_hook(&init_net, &hook3);
121     nf_unregister_net_hook(&init_net, &hook4);
122     nf_unregister_net_hook(&init_net, &hook5);
123     nf_unregister_net_hook(&init_net, &hook6);
124 }
125 }
126
```

将其编译，插入内核

```
[07/25/21]seed@VM:~/.../packet_filter$ sudo insmod seedFilter.ko
[07/25/21]seed@VM:~/.../packet_filter$ lsmod | grep seedFilter
seedFilter                16384  0
[07/25/21]seed@VM:~/.../packet_filter$ █
```



在 10.9.0.5 上 ping 10.9.0.1

```
root@465b288c116a:/# ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
64 bytes from 10.9.0.1: icmp_seq=1 ttl=64 time=0.073 ms
64 bytes from 10.9.0.1: icmp_seq=2 ttl=64 time=0.042 ms
64 bytes from 10.9.0.1: icmp_seq=3 ttl=64 time=0.048 ms
64 bytes from 10.9.0.1: icmp_seq=4 ttl=64 time=0.046 ms
64 bytes from 10.9.0.1: icmp_seq=5 ttl=64 time=0.046 ms
64 bytes from 10.9.0.1: icmp_seq=6 ttl=64 time=0.047 ms
64 bytes from 10.9.0.1: icmp_seq=7 ttl=64 time=0.049 ms
64 bytes from 10.9.0.1: icmp_seq=8 ttl=64 time=0.049 ms
64 bytes from 10.9.0.1: icmp_seq=9 ttl=64 time=0.048 ms
64 bytes from 10.9.0.1: icmp_seq=10 ttl=64 time=0.073 ms
```

dmesg 查看信息，可以看到相关的标识

```
[ 5062.317839] Registering filters.
[ 5088.318369] *** LOCAL_OUT
[ 5088.318371] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5088.318378] *** POST_ROUTING
[ 5088.318378] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5088.318386] *** PRE_ROUTING
[ 5088.318386] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5088.318387] *** LOCAL_IN
[ 5088.318387] 127.0.0.1 --> 127.0.0.53 (UDP)
[ 5088.318531] *** LOCAL_OUT
[ 5088.318532] 192.168.2.120 --> 192.168.2.1 (UDP)
[ 5088.318536] *** POST_ROUTING
[ 5088.318537] 192.168.2.120 --> 192.168.2.1 (UDP)
[ 5088.327961] *** PRE_ROUTING
[ 5088.327962] 192.168.2.1 --> 192.168.2.120 (UDP)
[ 5088.327970] *** LOCAL_IN
[ 5088.327971] 192.168.2.1 --> 192.168.2.120 (UDP)
```

在 10.9.0.5 上 ping 192.168.60.5

```
root@96037e435ddc:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.211 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.090 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.073 ms
^C
```

dmesg 查看信息

```
[ 5248.354798] *** PRE_ROUTING
[ 5248.354800] 10.9.0.5 --> 192.168.60.5 (ICMP)
[ 5248.354808] *** FORWARD
[ 5248.354808] 10.9.0.5 --> 192.168.60.5 (ICMP)
[ 5248.354822] *** POST_ROUTING
[ 5248.354823] 10.9.0.5 --> 192.168.60.5 (ICMP)
[ 5248.354855] *** PRE_ROUTING
[ 5248.354855] 10.9.0.5 --> 192.168.60.5 (ICMP)
[ 5248.354862] *** FORWARD
[ 5248.354863] 10.9.0.5 --> 192.168.60.5 (ICMP)
[ 5248.354864] *** POST_ROUTING
[ 5248.354864] 10.9.0.5 --> 192.168.60.5 (ICMP)
[ 5248.354866] *** FORWARD
[ 5248.354866] 10.9.0.5 --> 192.168.60.5 (ICMP)
```

```
[07/23/21]seed@VM:~/.../Labsetup$ ping 10.9.0.5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=64 time=0.042 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=64 time=0.041 ms
64 bytes from 10.9.0.5: icmp_seq=3 ttl=64 time=0.045 ms
64 bytes from 10.9.0.5: icmp_seq=4 ttl=64 time=0.041 ms
64 bytes from 10.9.0.5: icmp_seq=5 ttl=64 time=0.042 ms
```

归纳出五个 netfilter hooks 的位置：

NF\_INET\_PRE\_ROUTING: 进行完版本号, 校验和等检测, 进入网络层的数据包通过此点, 此处可以进行目的地址转换

NF\_INET\_LOCAL\_IN: 经路由查找后, 送往本机的通过此点, 此处可以进行 INPUT 过滤

NF\_INET\_FORWARD: 要转发的包通过此点, 此处可以进行 FORWARD 过滤

NF\_INET\_LOCAL\_OUT: 将要通过网络设备发出的包通过此点, 此处可以进行内置的源地址转换

NF\_INET\_POST\_ROUTING: 本机发出的包通过此点, 此处可以进行 OUTPUT 过滤

3. 要实现以下功能: (1) 阻止其他计算机 ping VM (2) 阻止其他计算机 telnet VM

在 netfilter 中实现两个不同的 hook 函数, telnetFilter 和 ICMPFilter

在 telnetFilter 函数中提取 dst 端口为 23 且所用协议为 TCP 协议的报文, 输出其 src 和 dst 的 ip 地址, 并将其放入 NF\_DROP 中, 将其他的放入 NF\_ACCEPT 中

在 ICMPFilter 函数中提取所用协议为 ICMP 协议的报文, 输出其 src 和 dst 的 ip 地址和 mac 地址, 并将其放入 NF\_DROP 中, 将其他的放入 NF\_ACCEPT 中

在 registerFilter 函数中, 定义两个 hook, 分别用 telnetFilter 和 ICMPFilter 函数注册, 并在 removeFilter 中写明这两个 hook 对应的移除语句

其具体代码如下:

```
1#include <linux/kernel.h>
2#include <linux/module.h>
3#include <linux/netfilter.h>
4#include <linux/netfilter_ipv4.h>
5#include <linux/ip.h>
6#include <linux/tcp.h>
7#include <linux/udp.h>
8#include <linux/if_ether.h>
9#include <linux/inet.h>
10
11static struct nf_hook_ops hook1, hook2;
12
13unsigned int telnetFilter(void *priv, struct sk_buff *skb,
14                          const struct nf_hook_state *state)
15{
16    struct iphdr *iph;
17    struct tcphdr *tcph;
18
19    iph = ip_hdr(skb);
20    tcph = (void *)iph + iph->ihl*4;
21
22    if(iph->protocol == IPPROTO_TCP && tcph->dest == htons(23))
23    {
24        printk(KERN_INFO "DROPPING TELNET PACKET FROM %d.%d.%d.%d\n", ((unsigned char *)&iph->saddr)[0], ((unsigned char *)&iph->saddr)[1], ((unsigned char *)&iph->saddr)[2], ((unsigned char *)&iph->saddr)[3]);
25        return NF_DROP;
26    }
27    else
28
29
30        return NF_ACCEPT;
31    }
32
33}
34
35unsigned int ICMPFilter(void *priv, struct sk_buff *skb,
36                        const struct nf_hook_state *state)
37{
38    struct ethhdr *mac_header = (struct ethhdr *)skb->mac_header;
39    struct iphdr *ip_header = (struct iphdr *)skb->network_header;
40
41    if(!skb)
42        return NF_ACCEPT;
43    if(ip_header->protocol == IPPROTO_ICMP)
44    {
45        printk(KERN_INFO "SRC_MAC: %pM\n", mac_header->h_source);
46        printk(KERN_INFO "DST_MAC: %pM\n", mac_header->h_dest);
47        printk(KERN_INFO "SRC_IP: %pI4\n", &ip_header->saddr);
48        printk(KERN_INFO "DST_IP: %pI4\n", &ip_header->daddr);
49
50        printk(KERN_INFO "the protocol ICMP(%d) is dropped...\n", IPPROTO_ICMP);
51        return NF_DROP;
52    }
53
54    return NF_ACCEPT;
55}
56
57int registerFilter(void) {
58    printk(KERN_INFO "Registering filters.\n");
```

```

59
60 hook1.hook = telnetFilter;
61 hook1.hooknum = NF_INET_LOCAL_IN;
62 hook1.pf = PF_INET;
63 hook1.priority = NF_IP_PRI_FIRST;
64 nf_register_net_hook(&init_net, &hook1);
65
66 hook2.hook = ICMPFilter;
67 hook2.hooknum = NF_INET_LOCAL_OUT;
68 hook2.pf = PF_INET;
69 hook2.priority = NF_IP_PRI_FIRST;
70 nf_register_net_hook(&init_net, &hook2);
71
72 return 0;
73 }
74
75 void removeFilter(void) {
76     printk(KERN_INFO "The filters are being removed.\n");
77     nf_unregister_net_hook(&init_net, &hook1);
78     nf_unregister_net_hook(&init_net, &hook2);
79 }
80
81 module_init(registerFilter);
82 module_exit(removeFilter);
83
84 MODULE_LICENSE("GPL");

```

在 docker 中进行测试

```

[07/23/21]seed@VM:~/.../Labsetup$ dockps
993d9a2666d8  seed-router
465b288c116a  hostA-10.9.0.5
7e615aa03d5f  host1-192.168.60.5
d2ad69c9095f  host3-192.168.60.7
5587b9c6812c  host2-192.168.60.6
[07/23/21]seed@VM:~/.../Labsetup$ docksh 46
root@465b288c116a:/# ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
^C
--- 10.9.0.1 ping statistics ---
9 packets transmitted, 0 received, 100% packet loss, time 8179ms

root@465b288c116a:/# dmesg
[ 0.000000] Linux version 5.4.0-54-generic (build@lcy01-amd64-024) (gcc vers
ion 9.3.0 (Ubuntu 9.3.0-17ubuntu1~20.04)) #60-Ubuntu SMP Fri Nov 6 10:37:59 UTC
2020 (Ubuntu 5.4.0-54.60-generic 5.4.65)

```

在 10.9.0.5 的 hostA 机 ping 10.9.0.1 时，所有包都被丢弃，在 dmesg 中可以看到丢包的 ip 地址和 mac 地址信息

```

[ 2188.283728] the protocol ICMP(1) is dropped...
[ 2189.304420] SRC_MAC:3f:00:00:9b:1c:4d
[ 2189.304421] DST_MAC:00:a4:81:00:00:8b
[ 2189.304422] SRC_IP:10.9.0.1
[ 2189.304422] DST_IP:10.9.0.5
[ 2189.304423] the protocol ICMP(1) is dropped...
[ 2190.332099] SRC_MAC:3f:00:00:9b:1c:4d
[ 2190.332100] DST_MAC:00:a4:81:00:00:8b
[ 2190.332100] SRC_IP:10.9.0.1
[ 2190.332100] DST_IP:10.9.0.5
[ 2190.332101] the protocol ICMP(1) is dropped...
[ 2191.354196] SRC_MAC:e3:00:00:c6:73:c3
[ 2191.354197] DST_MAC:00:a4:81:00:00:68
[ 2191.354197] SRC_IP:10.9.0.1
[ 2191.354198] DST_IP:10.9.0.5
[ 2191.354198] the protocol ICMP(1) is dropped...

```

在 telnet 10.9.0.1 时，telnet 失败（一直在 trying）

```

root@465b288c116a:/# telnet 10.9.0.1
Trying 10.9.0.1...
^C
-

```

在 dmesg 中可以看到一直在 dropping



```
[ 2283.612908] DROPPING TELNET PACKET FROM 10.9.0.5
[ 2284.633179] DROPPING TELNET PACKET FROM 10.9.0.5
[ 2286.652219] DROPPING TELNET PACKET FROM 10.9.0.5
[ 2290.713220] DROPPING TELNET PACKET FROM 10.9.0.5
[ 2298.907650] DROPPING TELNET PACKET FROM 10.9.0.5
```

故完成 task 的需求，防火墙试验成功

## 4 Task2: Experimenting with Stateless Firewall Rules

### 4.3 Task 2.A Protecting the Router

在路由器上输入以下命令设置 iptables,

```
root@993d9a2666d8:/# iptables -A INPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@993d9a2666d8:/# iptables -A OUTPUT -p icmp --icmp-type echo-request -j ACCEPT
root@993d9a2666d8:/# iptables -P OUTPUT DROP
root@993d9a2666d8:/# iptables -P INPUT DROP
root@993d9a2666d8:/#
```

通过 ifconfig 命令查看路由器的 ip 地址

```
root@993d9a2666d8:/# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.9.0.11 netmask 255.255.255.0 broadcast 10.9.0.255
    ether 02:42:0a:09:00:0b txqueuelen 0 (Ethernet)
    RX packets 83 bytes 8038 (8.0 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 50 bytes 3330 (3.3 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.60.11 netmask 255.255.255.0 broadcast 192.168.60.255
    ether 02:42:c0:a8:3c:0b txqueuelen 0 (Ethernet)
    RX packets 115 bytes 9920 (9.9 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 1 bytes 42 (42.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

在 10.9.0.5 的 HostA 上 ping 10.9.0.11 发现 ping 不通，通过 telnet 命令也无法远程登录

```
root@465b288c116a:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
^C
--- 10.9.0.11 ping statistics ---
9 packets transmitted, 0 received, 100% packet loss, time 8187ms

root@465b288c116a:/# telnet 10.9.0.11
Trying 10.9.0.11...
^C
```

解释每条 rules:

(1) `iptables -A INPUT -p icmp --icmp-type echo-reply -j ACCEPT`  
-A INPUT 表示向 INPUT 链中增加一条规则，-p icmp 表示指定协议为 icmp 协议，--icmp echo-reply 表示 icmp 类型为 echo-reply，-j ACCEPT 表示封包放行，即处理完之后不再比较其他规则。

该行规则表示向 INPUT 链中增加一条对 echo-reply 类型的 ICMP 规则，在处理后再不再比较其他规则封包放行。

(2) `iptables -A OUTPUT -p icmp --icmp-type echo-request -j ACCEPT`  
-A OUTPUT 表示向 OUTPUT 链中增加一条规则，-p icmp 表示指定协议为 icmp 协议，--

icmp echo-request 表示 icmp 类型为 echo-request, -j ACCEPT 表示封包放行, 即处理完之后不再比较其他规则。

该行规则表示向 OUTPUT 链中增加一条对 echo-request 类型的 ICMP 规则, 在处理后再不再比较其他规则封包放行。

(3) iptables -P OUTPUT DROP

-P 为设置链的默认策略, 该规则设置 OUTPUT 链的默认策略为丢包。

(4) iptables -P INPUT DROP

-P 为设置链的默认策略, 该规则设置 INPUT 链的默认策略为丢包。

在试验之后, 将 iptables 设置移除, 返回其初始设置, 并将 INPUT OUTPUT 默认策略改为 ACCEPT

```
root@993d9a2666d8:/# iptables -F
root@993d9a2666d8:/# iptables -P OUTPUT ACCEPT
root@993d9a2666d8:/# iptables -P INPUT ACCEPT
```

## 4.4 Task2.B: Protecting the Internal Network

为保护内网 192.168.60.0/24, 需要使用 FORWARD 链, 需要对 ICMP 流量做出如下限制:

- (1) 外部主机不能 ping 通内部主机
- (2) 外部主机可以 ping 通路由器
- (3) 内部主机可以 ping 通外部主机
- (4) 其他在内外网络之间的包应该被阻塞

首先使用 iptables -p icmp -h 查看相关信息

```
root@993d9a2666d8:/# iptables -p icmp -h
iptables v1.8.4

Usage: iptables -[ACD] chain rule-specification [options]
       iptables -I chain [rulenum] rule-specification [options]
       iptables -R chain rulenum rule-specification [options]
       iptables -D chain rulenum [options]
       iptables -[LS] [chain [rulenum]] [options]
       iptables -[FZ] [chain] [options]
       iptables -[NX] chain
       iptables -E old-chain-name new-chain-name
       iptables -P chain target [options]
       iptables -h (print this help information)

Commands:
Either long or short options are allowed.
--append -A chain          Append to chain
--check  -C chain          Check for the existence of a rule
--delete -D chain          Delete matching rule from chain
--delete -D chain rulenum  Delete rule rulenum (1 = first) from chain
--insert -I chain [rulenum] Insert in chain as rulenum (default 1=first)
--replace -R chain rulenum
```

在不做任何保护的时候, 内部外部之间可以相互 ping 通

```
root@7e615aa03d5f:/# ping 10.9.0.5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=63 time=0.064 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=63 time=0.059 ms
64 bytes from 10.9.0.5: icmp_seq=3 ttl=63 time=0.058 ms
64 bytes from 10.9.0.5: icmp_seq=4 ttl=63 time=0.093 ms
```



```
root@465b288c116a:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.144 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.157 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.069 ms
```

设定以下 iptables 规则，并设置 FORWARD 链的默认措施为 DROP

```
root@993d9a2666d8:/# iptables -A INPUT -p icmp -j ACCEPT
```

```
root@993d9a2666d8:/# iptables -A FORWARD -p icmp --icmp-type echo-request -i eth1
-o eth0 -j ACCEPT
```

```
root@993d9a2666d8:/# iptables -A FORWARD -p icmp --icmp-type echo-reply -d 192.168
.60.0/24 -j ACCEPT
```

```
root@993d9a2666d8:/# iptables -P FORWARD DROP
```

查看 filter 表的全部规则如下：

```
root@993d9a2666d8:/# iptables -t filter -L -n
```

Chain INPUT (policy ACCEPT)

target	prot	opt	source	destination
ACCEPT	icmp	--	0.0.0.0/0	0.0.0.0/0

Chain FORWARD (policy DROP)

target	prot	opt	source	destination	icmptype
ACCEPT	icmp	--	0.0.0.0/0	192.168.60.0/24	0
ACCEPT	icmp	--	0.0.0.0/0	0.0.0.0/0	8

Chain OUTPUT (policy ACCEPT)

target	prot	opt	source	destination
--------	------	-----	--------	-------------

验证这些规则是否满足需求

1 外部主机可以 ping 通路由器（两个端口）

```
root@465b288c116a:/# ping 192.168.60.11
PING 192.168.60.11 (192.168.60.11) 56(84) bytes of data.
64 bytes from 192.168.60.11: icmp_seq=1 ttl=64 time=0.069 ms
64 bytes from 192.168.60.11: icmp_seq=2 ttl=64 time=0.066 ms
64 bytes from 192.168.60.11: icmp_seq=3 ttl=64 time=0.076 ms
64 bytes from 192.168.60.11: icmp_seq=4 ttl=64 time=0.049 ms
```

```
root@465b288c116a:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.072 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.053 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.048 ms
```

2.外部主机无法 ping 通内部

```
root@465b288c116a:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
^C
--- 192.168.60.5 ping statistics ---
11 packets transmitted, 0 received, 100% packet loss, time 10257ms
```

3.路由器可以 ping 通外部主机

```
root@993d9a2666d8:/# ping 10.9.0.5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=64 time=0.097 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=64 time=0.049 ms
64 bytes from 10.9.0.5: icmp_seq=3 ttl=64 time=0.050 ms
```

4.内部主机可以 ping 通外部主机

```

root@7e615aa03d5f:/# ping 10.9.0.5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=63 time=0.064 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=63 time=0.061 ms
64 bytes from 10.9.0.5: icmp_seq=3 ttl=63 time=0.059 ms
64 bytes from 10.9.0.5: icmp_seq=4 ttl=63 time=0.059 ms

```

5.其他在内部主机和外部主机之间的流量被阻塞（如 telnet 不上）

```

root@7e615aa03d5f:/# telnet 10.9.0.5
Trying 10.9.0.5...
^C

```

## 4.5 Task2.C: Protecting Internet Servers

为了保护内部网的服务器（192.168.60.0/24 网段），需要满足以下需求：

- (1) 外部主机只能 telnet 192.168.60.5,无法 telnet 其他内网服务器
- (2) 外部主机无法到达内部服务器
- (3) 内部主机可以访问内部服务器
- (4) 内部主机不能访问外部服务器
- (5) 不能使用 connection tracking 机制

设定以下 iptables 规则，并设置 FORWARD 链的默认措施为 DROP

```

root@993d9a2666d8:/# iptables -A FORWARD -p tcp --sport 23 -d 192.168.60.5 -j ACCEPT
root@993d9a2666d8:/# iptables -A FORWARD -p tcp -s 192.168.60.0/24 -d 192.168.60.0/24 -j ACCEPT
root@993d9a2666d8:/# iptables -A FORWARD -p tcp --dport 23 -d 192.168.60.5 -j ACCEPT
root@993d9a2666d8:/# iptables -A FORWARD -p tcp -i eth1 -o eth0 --sport 23 -j ACCEPT

```

查看 filter 表的全部规则如下：

```

root@993d9a2666d8:/# iptables -t filter -L -n
Chain INPUT (policy ACCEPT)
target     prot opt source                destination
Chain FORWARD (policy DROP)
target     prot opt source                destination
ACCEPT     tcp  --  0.0.0.0/0              192.168.60.5           tcp spt:23
ACCEPT     tcp  --  192.168.60.0/24        192.168.60.0/24
ACCEPT     tcp  --  0.0.0.0/0              0.0.0.0/0              tcp spt:23
ACCEPT     tcp  --  0.0.0.0/0              192.168.60.5           tcp dpt:23
Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination
root@993d9a2666d8:/#

```

验证这些规则是否满足需求

- 1.内部主机可以 telnet 内部服务器，也可以 ping 通内部主机

```
root@7e615aa03d5f:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
5587b9c6812c login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage
```

```
seed@5587b9c6812c:~$ ping 192.168.60.6
PING 192.168.60.6 (192.168.60.6) 56(84) bytes of data.
64 bytes from 192.168.60.6: icmp_seq=1 ttl=64 time=70.7 ms
64 bytes from 192.168.60.6: icmp_seq=2 ttl=64 time=0.035 ms
64 bytes from 192.168.60.6: icmp_seq=3 ttl=64 time=0.032 ms
64 bytes from 192.168.60.6: icmp_seq=4 ttl=64 time=0.032 ms
```

## 2.外部主机无法 ping 通内部主机

```
root@465b288c116a:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
^C
--- 192.168.60.5 ping statistics ---
4 packets transmitted, 0 received, 100% packet loss, time 3056ms
```

## 3.内部主机无法 ping 通外部主机

```
seed@5587b9c6812c:~$ ping 10.9.0.5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
^C
--- 10.9.0.5 ping statistics ---
4 packets transmitted, 0 received, 100% packet loss, time 3445ms
```

## 4.外部主机可以 telnet 到 192.168.60.5

```
root@465b288c116a:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
seed
Ubuntu 20.04.1 LTS
seed
7e615aa03d5f login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage

This system has been minimized by removing packages and content that are
not required on a system that users do not log into.
```

## 5.外部主机无法 telnet 到其他内部服务器

```
root@465b288c116a:/# telnet 192.168.60.6
Trying 192.168.60.6...
^C
root@465b288c116a:/#
```



## 5 Task3: Connection Tracking and Stateful Firewall

### 5.1 Task 3.A: Experiment with the Connection Tracking

使用 conntrack -L 命令查看连接信息

```
root@993d9a2666d8:/# conntrack -L
conntrack v1.4.5 (conntrack-tools): 0 flow entries have been shown.
```

ICMP 实验: 在 10.9.0.5 上 ping 192.168.60.5

```
root@96037e435ddc:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.145 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.066 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.066 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.058 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.061 ms
64 bytes from 192.168.60.5: icmp_seq=6 ttl=63 time=0.059 ms
```

在路由器 conntrack -L 出现如下连接追踪信息, 状态维持时间为 29 秒

```
root@e93a0707dc73:/# conntrack -L
icmp      1 29 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=28 src=192.168.60.5
dst=10.9.0.5 type=0 code=0 id=28 mark=0 use=1
```

UDP 实验: 在 192.168.60.5 处输入以下命令

```
root@7e615aa03d5f:/# nc -lu 9090
boom
```

在 10.9.0.5 处输入以下命令

```
root@465b288c116a:/# nc -u 192.168.60.5 9090
boom
```

输入 boom,在回车之后, 建立连接, 在 192.168.60.5 出现输入信息

```
root@993d9a2666d8:/# conntrack -L
udp      17 21 src=10.9.0.5 dst=192.168.60.5 sport=54116 dport=9090 [UNREPLIED] sr
c=192.168.60.5 dst=10.9.0.5 sport=9090 dport=54116 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@993d9a2666d8:/#
```

在路由器 conntrack -L 出现如下连接追踪信息, 状态维持时间为 21 秒

TCP 实验: 在 192.168.60.5 处输入以下命令

```
root@7e615aa03d5f:/# nc -l 9090
boomboom
```

在 10.9.0.5 处输入以下命令

```
root@465b288c116a:/# nc 192.168.60.5 9090
boomboom
```

输入 boomboom,在回车之后, 建立连接, 在 192.168.60.5 出现输入信息

```
root@993d9a2666d8:/# conntrack -L
tcp      6 431964 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=53660 dport=9090
src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=53660 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@993d9a2666d8:/#
```

在路由器 conntrack -L 出现如下连接追踪信息，状态维持时间 431964 秒，时间明显长于 icmp 和 udp 连接

## 5.2 Task3.B: Setting Up a Stateful Firewall

在路由器上输入以下规则

```
root@993d9a2666d8:/# iptables -A FORWARD -p tcp -s 192.168.60.0/24 -j ACCEPT

root@993d9a2666d8:/# iptables -A FORWARD -p tcp -m conntrack -d 192.168.60.5 --dport 23 --ctstate ESTABLISHED,NEW -j ACCEPT

root@993d9a2666d8:/# iptables -A FORWARD -p tcp -m conntrack -d 192.168.60.0/24 --ctstate ESTABLISHED -j ACCEPT
```

查看 filter 表的全部规则如下：

```
root@993d9a2666d8:/# iptables -t filter -L -n
Chain INPUT (policy ACCEPT)
target     prot opt source                destination

Chain FORWARD (policy DROP)
target     prot opt source                destination
ACCEPT     tcp  --  0.0.0.0/0              192.168.60.0/24      ctstate ESTABLISHED
ACCEPT     tcp  --  0.0.0.0/0              192.168.60.5         ctstate NEW,ESTABLISHED tcp dpt:23
ACCEPT     tcp  --  192.168.60.0/24       0.0.0.0/0

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination
```

验证这些规则是否满足需求

1.内部主机可以 telnet 外部服务器，也与外部主机（10.9.0.5）建立连接

In->out

```
root@7e615aa03d5f:/# telnet 10.9.0.5
Trying 10.9.0.5...
Connected to 10.9.0.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
465b288c116a login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage
```

```
root@465b288c116a:/# nc -l 9090
123
```

```
root@7e615aa03d5f:/# nc 10.9.0.5 9090
123
```

2.内部主机可以 telnet 内部主机，也可以和内部主机建立连接

```
root@7e615aa03d5f:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
5587b9c6812c login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

* Documentation:  https://help.ubuntu.com
* Management:    https://landscape.canonical.com
* Support:       https://ubuntu.com/advantage
```

```

root@7e615aa03d5f:/# nc -nv 192.168.60.6 9090
Connection to 192.168.60.6 9090 port [tcp/*] succeeded!
12345
root@5587b9c6812c:/# nc -l 9090
12345

```

3.外部主机无法和内部主机建立连接

```

root@465b288c116a:/# nc 192.168.60.5 9090
1234
root@7e615aa03d5f:/# nc -l 9090

```

4.外部主机可以 telnet 192.168.60.5,且不能 telnet 其他内部主机 (192.168.60.6)

```

root@465b288c116a:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
7e615aa03d5f login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

This system has been minimized by removing packages and content that are
not required on a system that users do not log into.

root@465b288c116a:/# telnet 192.168.60.6
Trying 192.168.60.6...
^C
root@465b288c116a:/#

```

比较使用了连接跟踪(-m conntrack)和未使用该机制的方法，其中使用连接跟踪的方法写的规则更为简单，简化了规则设计并提高了效率，但是他需要更多的物理内存和存储空间

## 6 Task 4: Limiting Network Traffic

在路由器上运行以下命令

```

root@993d9a2666d8:/# iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute --limit-burst 5 -j ACCEPT
root@993d9a2666d8:/# iptables -A FORWARD -s 10.9.0.5 -j DROP

```

在 10.9.0.5 上 ping 192.168.60.5

```

root@465b288c116a:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.112 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.061 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.060 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.094 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.062 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.059 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.058 ms

```

可以直观的看到输出增长比较慢，即接收到的包的频率较慢，实现了对流量的限制

```

root@993d9a2666d8:/# iptables -D FORWARD -s 10.9.0.5 -j DROP

```

去掉第二行规则后重复发送



```

root@465b288c116a:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.064 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.062 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.066 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.061 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.066 ms
64 bytes from 192.168.60.5: icmp_seq=6 ttl=63 time=0.072 ms

```

直观的看到与平常 ping 的增长速度相同，即接收到的包的频率和之前相同，没有实现流量的限制

对比下表表明第二行命令是需要的，否则起不到流量限制的效果，其原因为默认规则为 ACCEPT 下，不满足的其他的报文也会被 accept，故所有流量都得不到限制

## 7 Task 5: Load Balancing

在 192.168.60.5, 192.168.60.6, 192.168.60.7 上分别输入以下命令，

```

root@7e615aa03d5f:/# nc -luk 8080

```

```

root@5587b9c6812c:/# nc -luk 8080

```

```

root@d2ad69c9095f:/# nc -luk 8080

```

1.使用 nth mode:

在路由器中输入以下规则，即每三个报文的第一个发送到 192.168.60.5: 8080 端口，在发送后，剩下的报文中每两个报文的第一个发送到 192.168.60.7: 8080 端口，剩下的所有报文发送到 192.168.60.6: 8080 端口，以起到负载均衡的效果

```

root@993d9a2666d8:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 3
--packet 0 -j DNAT --to-destination 192.168.60.5:8080

```

```

root@993d9a2666d8:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 2
--packet 0 -j DNAT --to-destination 192.168.60.7:8080

```

```

root@993d9a2666d8:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -j DNAT --to-destination 192.168.60.6:8080

```

在 10.9.0.5 上输入以下命令

```

root@465b288c116a:/# echo hello1 | nc -u 10.9.0.11 8080
^C
root@465b288c116a:/# echo hello2 | nc -u 10.9.0.11 8080
^C
root@465b288c116a:/# echo hello3 | nc -u 10.9.0.11 8080
^C

```

可见 hello1 被发送到 192.168.60.5: 8080 端口

```

root@7e615aa03d5f:/# nc -luk 8080
hello1

```

hello3 被发送到 192.168.60.6: 8080 端口

```

root@5587b9c6812c:/# nc -luk 8080
hello3

```

hello2 被发送到 192.168.60.7: 8080 端口

```

root@d2ad69c9095f:/# nc -luk 8080
hello2

```

实现了负载均衡

## 2.使用 random mode

首先清除之前的 iptables 规则，路由器中输入以下规则，即将以 0.33 的概率将报文发到 192.168.60.5: 8080 端口，在发送后，剩下的报文中将报文以 0.5 的概率发送到 192.168.60.6: 8080 端口，剩下的所有报文发送到 192.168.60.7: 8080 端口，以起到负载均衡的效果

```
root@993d9a2666d8:/# iptables -F
root@993d9a2666d8:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.33 -j DNAT --to-destination 192.168.60.5:8080
root@993d9a2666d8:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.5 -j DNAT --to-destination 192.168.60.6:8080
root@993d9a2666d8:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -j DNAT --to-destination 192.168.60.6:8080
root@993d9a2666d8:/#
```

在 10.9.0.5 上输入以下命令

```
root@465b288c116a:/# echo hello1 | nc -u 10.9.0.11 8080
^C
root@465b288c116a:/# echo hello2 | nc -u 10.9.0.11 8080
^C
root@465b288c116a:/# echo hello3 | nc -u 10.9.0.11 8080
^C
root@465b288c116a:/# echo hello4 | nc -u 10.9.0.11 8080
^C
root@465b288c116a:/# echo hello5 | nc -u 10.9.0.11 8080
^C
root@465b288c116a:/# echo hello6 | nc -u 10.9.0.11 8080
^C
root@465b288c116a:/# echo hello7 | nc -u 10.9.0.11 8080
^C
```

可见 hello1, hello4, hello7 被发送到 192.168.60.5: 8080 端口

```
root@7e615aa03d5f:/# nc -luk 8080
hello1
hello4
hello7
█
```

Hello3, hello6 被发送到 192.168.60.6: 8080 端口

```
root@5587b9c6812c:/# nc -luk 8080
hello3
hello6
█
```

Hello2, hello5 被发送到 192.168.60.7: 8080 端口

```
root@d2ad69c9095f:/# nc -luk 8080
hello2
hello5
█
```

实现了负载均衡